

## Importing the Dependencies

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.m

```
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

## Data Collection and Processing

```
# loading the csv data to a Pandas DataFrame
heart_data = pd.read_csv('/content/drive/MyDrive/heart.csv')
```

```
# print first 5 rows of the dataset
heart_data.head()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca
0	63	1	3	145	233	1	0	150	0	2.3	0	0
1	37	1	2	130	250	0	1	187	0	3.5	0	0
2	41	0	1	130	204	0	0	172	0	1.4	2	0
3	56	1	1	120	236	0	1	178	0	0.8	2	0
4	57	0	0	120	354	0	1	163	1	0.6	2	0

```
# print last 5 rows of the dataset
heart_data.tail()
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	c
298	57	0	0	140	241	0	1	123	1	0.2	1	1
299	45	1	3	110	264	0	1	132	0	1.2	1	1
300	68	1	0	144	193	1	1	141	0	3.4	1	1
301	57	1	0	130	131	0	1	115	1	1.2	1	1
302	57	0	1	130	236	0	0	174	0	0.0	1	1

```
# number of rows and columns in the dataset
```

```
heart_data.shape
```

```
(303, 14)
```

```
# getting some info about the data
```

```
heart_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   age         303 non-null   int64  
 1   sex         303 non-null   int64  
 2   cp          303 non-null   int64  
 3   trestbps    303 non-null   int64  
 4   chol        303 non-null   int64  
 5   fbs         303 non-null   int64  
 6   restecg     303 non-null   int64  
 7   thalach     303 non-null   int64  
 8   exang       303 non-null   int64  
 9   oldpeak     303 non-null   float64 
10   slope       303 non-null   int64  
11   ca          303 non-null   int64  
12   thal        303 non-null   int64  
13   target      303 non-null   int64  
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

```
# checking for missing values
```

```
heart_data.isnull().sum()
```

```
age         0
sex         0
cp          0
trestbps    0
chol        0
fbs         0
restecg     0
thalach     0
exang       0
oldpeak     0
slope       0
ca          0
thal        0
target      0
dtype: int64
```

```
# statistical measures about the data
```

```
heart_data.describe()
```

	age	sex	cp	trestbps	chol	fbs	res
<b>count</b>	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.00
<b>mean</b>	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	0.52
<b>std</b>	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	0.52
<b>min</b>	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.00
<b>25%</b>	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	0.00

```
# checking the distribution of Target Variable
heart_data['target'].value_counts()
```

```
1    165
0    138
Name: target, dtype: int64
```

1 --> Defective Heart

0 --> Healthy Heart

### Splitting the Features and Target

```
X = heart_data.drop(columns='target', axis=1)
Y = heart_data['target']
```

```
print(X)
```

```

   age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  \
0    63   1   3    145    233   1         0     150      0      2.3
1    37   1   2    130    250   0         1     187      0      3.5
2    41   0   1    130    204   0         0     172      0      1.4
3    56   1   1    120    236   0         1     178      0      0.8
4    57   0   0    120    354   0         1     163      1      0.6
..  ...  ...  ..  ...    ...  ...  ...    ...    ...  ...
298  57   0   0    140    241   0         1     123      1      0.2
299  45   1   3    110    264   0         1     132      0      1.2
300  68   1   0    144    193   1         1     141      0      3.4
301  57   1   0    130    131   0         1     115      1      1.2
302  57   0   1    130    236   0         0     174      0      0.0

   slope  ca  thal
0       0   0     1
1       0   0     2
2       2   0     2
3       2   0     2
4       2   0     2
..  ...  ..  ...
298    1   0     3
299    1   0     3
300    1   2     3
301    1   1     3
302    1   1     2
```

```
[303 rows x 13 columns]
print(Y)

0      1
1      1
2      1
3      1
4      1
..
298    0
299    0
300    0
301    0
302    0
Name: target, Length: 303, dtype: int64
```

## Splitting the Data into Training data & Test Data

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, stratify=Y, random_state=1)

print(X.shape, X_train.shape, X_test.shape)

(303, 13) (242, 13) (61, 13)
```

## Model Training

### Logistic Regression

```
model = LogisticRegression()

# training the LogisticRegression model with Training data
model.fit(X_train, Y_train)

/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:818: ConvergenceWarning:
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
LogisticRegression()
```

## Model Evaluation

### Accuracy Score

```
# accuracy on training data
X_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)

print('Accuracy on Training data : ', training_data_accuracy)

    Accuracy on Training data :    0.8512396694214877

# accuracy on test data
X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)

print('Accuracy on Test data : ', test_data_accuracy)

    Accuracy on Test data :    0.819672131147541
```

## Building a Predictive System

```
input_data = (62,0,0,140,268,0,0,160,0,3.6,0,2,2)


# change the input data to a numpy array
input_data_as_numpy_array= np.asarray(input_data)

# reshape the numpy array as we are predicting for only on instance
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)

prediction = model.predict(input_data_reshaped)
print(prediction)

if (prediction[0]== 0):
    print('The Person does not have a Heart Disease')
else:
    print('The Person has Heart Disease')

[0]
The Person does not have a Heart Disease
/usr/local/lib/python3.7/dist-packages/sklearn/base.py:451: UserWarning: X does not have valid feature names, but"
```



## Saving the trained model

```
import pickle

filename = 'heart_disease_model.sav'
pickle.dump(model, open(filename, 'wb'))

# loading the saved model
loaded_model = pickle.load(open('heart disease model.sav', 'rb'))
```

```
for column in X.columns:  
    print(column)
```

```
↳ age  
sex  
cp  
trestbps  
chol  
fbs  
restecg  
thalach  
exang  
oldpeak  
slope  
ca  
thal
```